

UNIVERSITÀ DEGLI STUDI DI MILANO

Dipartimento di Informatica

Architecting an End-to-End Secure Event-Driven Pipeline

*A Comprehensive Implementation of Kafka KRaft, Mutual
TLS, and Chaos-Tested Fault Tolerance in Go*

Candidate: Saji Hamati

Matriculation No: 55364A

Email: saji.hamati@studenti.unimi.it

Academic Year: 2025/2026

March 31, 2026

Abstract

As modern enterprise systems transition from monolithic architectures to distributed microservices, the complexity of maintaining data integrity, system availability, and network security increases exponentially. This technical report details the design, implementation, and rigorous stress-testing of a highly scalable, fault-tolerant Order Processing System. Built entirely using the Go programming language, the system utilizes a ZooKeeper-less Apache Kafka cluster (KRaft mode) as its central event broker, and a containerized PostgreSQL database for ACID-compliant persistence. A primary focus of this research is the implementation of a "Zero Trust" security model enforcing Mutual Transport Layer Security (mTLS). Furthermore, this report details the mathematical load-balancing strategies employed for parallel processing and documents a "Chaos Engineering" experiment, wherein active cluster nodes were intentionally terminated under heavy synthetic load to empirically validate the architecture's fault tolerance and zero-data-loss guarantees.

Contents

1	Introduction and Architectural Context	3
1.1	Language Selection: Go (Golang)	3
2	The Message Broker: Apache Kafka in KRaft Mode	3
2.1	Role Separation: Dedicated Controllers and Brokers	3
2.2	The Raft Consensus Algorithm in Action	4
3	Military-Grade Security: Implementing mTLS	4
3.1	Public Key Infrastructure (PKI) Lifecycle	4
3.2	The Mutual TLS (mTLS) Handshake	4
4	Database Hardening and The "Permission Trap"	5
4.1	The Filesystem Permission Constraint	5
4.2	The Container Orchestration Solution	5
5	Load Balancing and Resource Control	5
5.1	Mathematical Hash Partitioning	6
6	Observability and System Visualization	6
6.1	Cluster Monitoring: Kafka UI	6
6.2	Database Management: pgAdmin 4	6
7	Performance Evaluation and Chaos Engineering	7
7.1	The Asynchronous Bash Script	7
7.2	Chaos Engineering: Active Node Termination	8
7.2.1	Observed System Behavior	8
8	Conclusion	8

1 Introduction and Architectural Context

The shift toward Event-Driven Architecture (EDA) allows independent services to scale dynamically but introduces new points of failure. The goal of this project was to design a system capable of ingesting high volumes of e-commerce orders, persisting them safely, and routing them to consumer workers without data loss or security vulnerabilities.

To achieve this, the architecture was segmented into three distinct layers:

1. **The Ingestion Layer:** A Producer REST API written in Go, responsible for payload validation and synchronous database commits.
2. **The Streaming Layer:** A 3-node Apache Kafka cluster responsible for holding the distributed, immutable log of events.
3. **The Processing Layer:** A Go Consumer worker that subscribes to the Kafka topic and processes the state transitions of the orders.

1.1 Language Selection: Go (Golang)

Go was selected as the primary language due to its native concurrency model based on Communicating Sequential Processes (CSP). By utilizing lightweight *Goroutines*, the Producer API can handle thousands of simultaneous HTTP connections with an incredibly small memory footprint. Furthermore, Go's standard library provides robust, low-level control over the cryptographic `crypto/tls` stack, which was critical for implementing custom Certificate Authority workflows.

2 The Message Broker: Apache Kafka in KRaft Mode

Traditionally, Kafka relied on Apache ZooKeeper to manage cluster metadata, leader elections, and broker discovery. This project utilizes Kafka's modern **KRaft** (**K**afka **R**aft) architecture, eliminating ZooKeeper entirely and consolidating the control plane.

2.1 Role Separation: Dedicated Controllers and Brokers

In our architecture, the control plane and data plane are strictly isolated into dedicated, independent containers to ensure resource isolation and enhanced fault tolerance. Rather than running nodes in a dual capacity, the cluster roles are physically separated:

- **The Control Plane (KRaft Controllers):** A dedicated quorum of controller nodes handles the Raft consensus. They manage the metadata of the cluster, track which brokers are online, and dictate which node acts as the leader for a given partition. Because they do not store physical event data, they are immune to disk-bound resource starvation.

- **The Data Plane (Brokers):** A separate group of broker nodes (e.g., `broker-1`, `broker-2`, `broker-3`) exclusively handles the physical I/O operations. They are solely responsible for writing the event streams to disk and serving fetch requests from the Go Consumer.

2.2 The Raft Consensus Algorithm in Action

The KRaft mode utilizes an internal metadata topic. When the Go application requests the creation of the `orders-topic`, the request is routed to the Active Controller. The Controller writes this state change to its local log and broadcasts it. Only when a majority of the nodes (2 out of 3) acknowledge the write does the cluster consider the topic officially created. This strict quorum logic prevents "split-brain" scenarios during network partitions.

3 Military-Grade Security: Implementing mTLS

In a standard cloud environment, the internal network cannot be inherently trusted. We implemented a Zero Trust architecture requiring all services to mathematically prove their identity before transmitting data.

3.1 Public Key Infrastructure (PKI) Lifecycle

Using OpenSSL, we established a localized Certificate Authority (CA). The PKI lifecycle consisted of:

1. Generating a Master Private Key and a self-signed Root CA Certificate (`ca.crt`).
2. Generating individual Private Keys for the Kafka Brokers, PostgreSQL, and Go Clients.
3. Creating Certificate Signing Requests (CSRs) for each entity, including Subject Alternative Names (SANs) to bind certificates to specific Docker hostnames.
4. Signing the CSRs with the Root CA, yielding fully authenticated x.509 certificates.

3.2 The Mutual TLS (mTLS) Handshake

Our Kafka implementation requires **Mutual Authentication**. When the Go Producer attempts to write an order, the following cryptographic sequence occurs:

1. The Go Client connects to Kafka over TCP port 19092.
2. Kafka presents its Broker Certificate. Go validates the signature using the loaded `ca.crt`.
3. Kafka demands a client certificate. Go transmits `go-client.crt`.

4. Kafka inspects the client certificate. If not signed by the shared CA, the connection is dropped.
5. Upon mutual validation, the nodes perform a Diffie-Hellman key exchange, generating a symmetric AES-256 key to encrypt all subsequent data.

4 Database Hardening and The "Permission Trap"

PostgreSQL 15 serves as the source of truth for the system. We implemented **1-Way TLS** for the database connection, securing the transport layer to protect the credentials.

4.1 The Filesystem Permission Constraint

Securing PostgreSQL in Docker introduced a significant infrastructure challenge. The PostgreSQL engine enforces strict UNIX permission checks on private keys. If the `postgres.key` file possesses "world-readable" permissions (e.g., `0644`), the database will panic and refuse to boot. Because Docker volumes mapped from host machines do not properly translate UNIX permissions, direct mounting resulted in continuous fatal crashes.

4.2 The Container Orchestration Solution

To bypass this without custom Dockerfiles, an interceptive bootstrap script was injected into the `docker-compose.yml`.

```
1  command: >
2      sh -c "cp /tmp/secrets/postgres.* /var/lib/postgresql/ &&
3      chown postgres:postgres /var/lib/postgresql/postgres.* &&
4      chmod 0600 /var/lib/postgresql/postgres.key &&
5      exec docker-entrypoint.sh postgres -c ssl=on ..."
```

Listing 1: PostgreSQL Pre-boot Security Script

This script stages the certificates, copies them to the restricted `/var/lib/postgresql/` directory, assigns ownership to the `postgres` UNIX user, and explicitly locks the private key to `0600` (Read/Write for owner only).

5 Load Balancing and Resource Control

To handle massive throughput, the Kafka topic was divided into **3 Partitions**, allowing three Consumer Go-routines to process messages entirely in parallel. However, parallel processing introduces the risk of race conditions if updates for the *same* order are processed out of sequence (e.g., a "Shipped" event being processed before a "Pending" event).

5.1 Mathematical Hash Partitioning

To guarantee chronological ordering without sacrificing horizontal scalability, a **Keyed Hash Balancer** was implemented in the Go Producer. Instead of assigning messages to partitions randomly or via Round-Robin, the routing algorithm is defined as:

$$P_{target} = \text{CRC32}(\text{MessageKey}) \pmod{N}$$

where P_{target} is the destination partition, MessageKey is the unique Order ID (e.g., ORD-123), and N is the total number of partitions (3).

Because the Cyclic Redundancy Check (CRC32) hash of a specific string is deterministic, all events related to a specific order ID will always be routed to the exact same partition. Since Kafka guarantees strict message ordering *within a single partition*, the Consumer worker is mathematically guaranteed to process an individual customer's state transitions in the exact order they were received, while the aggregate system load remains perfectly balanced across all three brokers.

6 Observability and System Visualization

To maintain operational oversight of the distributed pipeline, dedicated graphical user interfaces (GUIs) were integrated into the container orchestration stack. In a highly decoupled architecture, these visualization tools proved essential for debugging mTLS connections and validating data ingestion during the stress-testing phases.

6.1 Cluster Monitoring: Kafka UI

Managing a distributed event broker via the Command Line Interface (CLI) can be opaque, especially when tracking metadata across a KRaft quorum. To mitigate this, **Kafka UI** was deployed as a lightweight container alongside the cluster. This interface provided real-time, visual telemetry on:

- **Cluster Health:** Monitoring the live status of the Active Controller and the synchronized Broker nodes.
- **Topic Metrics:** Visualizing the partition distribution, replication status, and leader assignments of the `orders-topic`.
- **Consumer Lag:** Tracking the offset delta between the Go Producer's ingestion rate and the Consumer Worker's processing speed, ensuring the system was not bottlenecking during high-load experiments.

6.2 Database Management: pgAdmin 4

To visualize the persistence layer and verify data integrity, **pgAdmin 4** was deployed as the administrative GUI for PostgreSQL. Because the database was strictly locked

down requiring 1-Way TLS, configuring the UI required passing the custom Root CA certificate (`ca.crt`) directly into the pgAdmin container via Docker volumes so the UI could mathematically verify the database server.

The web interface was exposed locally and secured with the following administrative credentials:

- **Email ID:** admin@admin.com
- **Password:** password123

Once authenticated, this UI allowed for the direct execution of SQL queries to verify that the ACID properties of the `orders` table were maintained during high-concurrency ingestion. Furthermore, it allowed us to query the `pg_stat_ssl` system view to definitively prove that the UI and the Go services were communicating with the database over an encrypted tunnel.

7 Performance Evaluation and Chaos Engineering

To validate the architecture's load balancing and fault tolerance, a custom testing apparatus was designed. Relying on static tools like Apache Bench was insufficient due to PostgreSQL's primary key constraints; thus, a dynamic injection script was required.

7.1 The Asynchronous Bash Script

A stress test was executed from the host machine targeting the Go Producer API. The script spawned thousands of asynchronous background processes utilizing the `bash &` operator. By generating unique UUIDs via the `$RANDOM` variable, the script simulated highly concurrent, distinct user traffic.

```
1 for i in {1..5000}; do
2   curl -s -X POST http://localhost:8080/orders \
3     -H "Content-Type: application/json" \
4     -d "{\"id\": \"$RANDOM-$i\", \"item\": \"Keyboard\", \"quantity\": 1}" \
5     > /dev/null &
6 done
7 wait
8 echo "Stress Test Complete!"
```

Listing 2: Asynchronous Stress Test Injection

During uninterrupted execution, the Go API effortlessly processed ~1,500 requests per second. The CPU overhead for the AES-256 mTLS encryption peaked at only 4% above plaintext baseline tests, proving the efficiency of Go's TLS implementation.

7.2 Chaos Engineering: Active Node Termination

To empirically test the system’s fault tolerance, a ”Chaos Engineering” experiment was conducted. The Kafka topic was configured with a **Replication Factor of 3** and the Producer was configured with `RequiredAcks: kafka.RequireAll`.

During the execution of the Bash Bomber script (at peak ingestion load), a critical cluster node was intentionally destroyed via the Docker CLI:

```
1 docker stop broker-2
```

7.2.1 Observed System Behavior

1. **Controller Detection:** The KRaft Active Controller immediately detected the loss of `broker-2` via missed heartbeat signals. 2. **Leader Election:** For any partitions where `broker-2` was the Leader, the Controller instantly promoted one of the remaining In-Sync Replicas (ISR) on `broker-1` or `broker-3` to the new Leader. 3. **Producer Buffering:** The Go Producer API detected the TCP connection drop. Instead of returning HTTP 500 errors to the client, the `kafka-go` library temporarily buffered the incoming HTTP requests in memory, requested the updated cluster metadata, and established a new mTLS tunnel with the newly elected partition leaders. 4. **Zero Data Loss Validation:** The system experienced a brief latency spike of approximately 1.5 seconds during the election phase. However, upon completion of the script, PostgreSQL and pgAdmin 4 confirmed exactly 5,000 unique rows had been successfully inserted and transitioned to the `PROCESSED` state. Zero HTTP requests were dropped, and zero duplicate records were created.

8 Conclusion

The successful design and deployment of this system validates that high-performance, containerized microservices can operate with uncompromising security and reliability. By migrating to Kafka KRaft, establishing a private PKI for Mutual TLS, overcoming strict database filesystem constraints, mathematically balancing parallel workloads, and surviving active node termination under heavy load, this architecture serves as a robust blueprint for modern, event-driven enterprise applications.